

# EEEM030 Group Assignment 2: Speech Recognition

*Assessment: 15% of module*

## 1. Overview

The second coursework assignment for *EEEM030 Speech & audio processing & recognition* is designed to give you an opportunity to experience the machine learning methodology in a small development team, and to practice the key algorithms that are used to extract features and to initialize, train and test your models. In doing so, you will get to apply the efficient recursive procedures to observe the effects of training using maximum likelihood, and to perform recognition for a simple, isolated-word recognition task.

The aim of the coursework is to develop a simple speech recognizer in Matlab that uses HMMs with a small vocabulary of eleven key words.

This is a group project. You will be assigned to a group with 2-6 members. Your group will submit **one joint technical report** (max. 5,000 words, plus references and any appendices). The report will include a distinct chapter from each team member, plus other essential parts submitted as a collective contribution (i.e., the abstract, introduction/background, conclusion and references). The whole document must be consistently formatted in a single PDF file that is less than 12MB. It must be legible and must not exceed 50 pages in total. Your final mark will assess both your **individual chapter** and the quality of the technical report as a whole.

The spokesperson for your group must submit the report as a **PDF file** to the EEEM030 assignment folder in SurreyLearn by **4pm Friday 6th December 2024** (week 11). Other team members can submit the **same document** if they wish, but only the spokesperson's submission will be marked.

The recommended structure for your group technical report is as follows, where each team member is to contribute **one** of the main central chapters highlighted in bold:

- Front page with university crest, module code (EEEM030), name of the assignment, title of the report, list of co-authors and copyright notice
- Abstract
- Table of Contents
- Introduction
- **First contribution chapter**
- **Second contribution chapter**
- ...
- **Last contribution chapter**
- Summary/conclusion
- References/bibliography
- Appendix/appendices

You may use publicly-available source code, where it is relevant, but any code you use (or adapt) that was written by someone outside your team **must be cited** with due reference (to avoid accusations of plagiarism), stating how it has been used. You will need to ensure you understand what it does and how it works in order to complete the assignment successfully, as there can be significant variations (e.g., around pdf normalisation or definition of null states).

The same principle applies within each individual contribution chapter, i.e., you may use code written by other team members in your work provided that you **attribute their contribution**. That means that you need to cross reference their work in your chapter. So make sure you understand how you'll do that when they provide access to their code: whether it is described in their chapter, included in the report as an appendix, or referenced externally in a public code repository (e.g., as a Github project).

As with all formally assessed coursework, you may discuss the concepts associated with the coursework with your peers in other teams, but not the details of any solution that you implement. You cannot share code between different teams. In line with University policy, you may use other sources, such as textbooks, lecture notes, articles, online tutorials and code libraries, but failure to cite them correctly may be viewed as plagiarism and trigger an academic misconduct investigation. So please reference all your sources carefully and thoroughly!

## 2. Model prototype and speech data

You are asked to initialize, train and test a set of hidden Markov models (HMMs), with one for each of the key words in the vocabulary. Each model is an 8-state HMM with 13-dimensional continuous probability density function (pdf), i.e.,  $N=8$  and  $K=13$ . The HMM has a strict left-right topology, as illustrated with some typical values of the state-transition probabilities  $A$  in Table 1. The output probabilities  $B$  are obtained with a 13-D multivariate Gaussian pdf that has a mean  $\mu$  and a diagonal covariance matrix  $\Sigma$ .

|   |     |     |     |     |     |     |     |     |     |
|---|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 0 | 1   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   |
| 0 | 0.8 | 0.2 | 0   | 0   | 0   | 0   | 0   | 0   | 0   |
| 0 | 0   | 0.8 | 0.2 | 0   | 0   | 0   | 0   | 0   | 0   |
| 0 | 0   | 0   | 0.8 | 0.2 | 0   | 0   | 0   | 0   | 0   |
| 0 | 0   | 0   | 0   | 0.8 | 0.2 | 0   | 0   | 0   | 0   |
| 0 | 0   | 0   | 0   | 0   | 0.8 | 0.2 | 0   | 0   | 0   |
| 0 | 0   | 0   | 0   | 0   | 0   | 0.8 | 0.2 | 0   | 0   |
| 0 | 0   | 0   | 0   | 0   | 0   | 0   | 0.8 | 0.2 | 0   |
| 0 | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0.8 | 0.2 |
| 0 | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   |

Table 1: State-transition probability matrix,  $A = \{\pi_i, a_{ij}, \eta_i\}$ , including entry and exit transitions.

The speech data comprise a collection of audio recordings from a variety of speakers. The vocabulary includes the set of key words listed in Table 2. Each audio file contains only one word. The filename indicates the index of the word spoken in that file and spells out the corresponding word.

You are initially provided with development data for training, and will later obtain a set of validation data to test your system.

| Index | 1      | 2     | 3      | 4     | 5      | 6     | 7     | 8       | 9      | 10      | 11      |
|-------|--------|-------|--------|-------|--------|-------|-------|---------|--------|---------|---------|
| Word  | "heed" | "hid" | "head" | "had" | "hard" | "hud" | "hod" | "hoard" | "hood" | "who'd" | "heard" |

Table 2: Set of isolated words making up the vocabulary for this recognition task.

## 3. Feature extraction, model initialisation, training and decoding

For feature extraction, use 13 Mel-frequency cepstral coefficients including the zeroth coefficient. Do not include the velocities and accelerations, i.e., the delta features and delta-delta-features. This will provide a 13-dimensional feature vector for each frame of audio data. The frame should take a 30ms block of audio samples, apply a tapered window and a hop size of 10ms, i.e., with a 20ms overlap.

For model initialisation, you are advised to compute the global mean and variance across the whole development set. For the mean, this should give you a 13-D vector of numbers. For the variance, we also want a set of 13 values in order to populate the diagonal of the  $13 \times 13$  covariance matrix. So you can either compute the variance for each dimension separately, or compute the full covariance and set the off-diagonal values equal to zero. You will use a scaled version of this global variance to set a variance floor in training, i.e., to limit how narrow these distributions can get.

From the number of frames extracted from the audio files, you can compute the average duration of each state, in frames, and hence an estimate for the self-loop transition probability  $a_{ii}$  from the average duration  $\tau$  as follows:

$$a_{ii} = \exp(-1 / (\tau - 1)).$$

Hence, you can obtain the probability of transition to the next state (all the transition options from a state must sum to one). These two values can be used to initialize all the self and onward state transitions in turn, as depicted for the illustrative example in Table 1, with the self-loop at 0.8 or four fifths, and the onward transition at 0.2 or one fifth.

For training, you must use the Baum-Welch equations to re-estimate new models, as described in detail in your lecture notes. This requires the implementation of the forward and backward procedures to obtain the forward and backward likelihoods,  $\alpha$  and  $\beta$  respectively, followed by the calculation of occupation likelihoods,  $\gamma$ , and transition likelihoods,  $\xi$ . These then increment the accumulators as each training file is processed, and finally provide the updated parameter values from one complete training epoch, where  $\hat{A} = \{\hat{a}_{ij}\}$  and  $\hat{B} = \{\hat{\mu}_i, \hat{\Sigma}_i\}$ .

For monitoring the training process on the development data, and later for evaluating the trained models with the validation data, you must use the Viterbi algorithm in the decoder. This is used to perform recognition by identifying the output label with the best match for a given test file, which allows you to calculate the recognition error rate and for a confusion matrix amongst word labels.

Note that the likelihood calculations may be computed directly as probabilities or, more efficiently, in the form of log-probs by taking the natural logarithm of each of the probability values. Importantly, to maintain the probability estimates within Matlab's range of numerical precision, you will need to re-scale the probabilities. This can either be done each time frame, for example, or using a global scale factor. You should expect to encounter this problem with this assignment, so inspect the probabilities obtained in your calculations and check for overflow/underflow in the values. Matlab may use infinity (Inf) or not-a-number (NaN) to denote such a result. I am happy to discuss solutions for this issue, and share this know-how with the cohort, so early questions on this are welcome, i.e., up to week 10.

## 4. Task

This assignment tackles an isolated word recognition (IWR) task by training HMMs according to the Expectation-Maximization method with the Baum-Welch equations, and running the recognizer by decoding with the Viterbi algorithm. You are provided development data for training and some test data for evaluation. You can record additional speech data to show further testing.

The activities that your group is asked to undertake are as follows:

1. Extract MFCC acoustic features
  - a. Compute features from the training data in the development set
  - b. Compute features from the test data in the evaluation set
  - c. Compute features from further test data using your own recordings
2. Initialise a set of prototype HMMs for each word in the vocabulary
  - a. Calculate global mean and variance statistics of the 13-dimensional feature data in the development set
  - b. Apply these values as flat-start prototype model parameters to all template models
3. Train the HMMs with the training data
  - a. Write implementations of the forward and backward procedures to calculate the forward and backward likelihoods with your feature data and models, according to the methods defined in your lectures
  - b. Investigate solutions to the underflow problem
  - c. Determine the occupation likelihoods to increment the accumulators for  $\hat{B}$
  - d. Determine the transition likelihoods to increment the accumulators for  $\hat{A}$

- e. Re-estimate the model parameters at the end of one iteration over the whole training data, and save the models  $\hat{\lambda} = \{\hat{A}, \hat{B}\}$
4. Evaluate the recognizer at the at each iteration on the development data
  - a. Write an implementation of the Viterbi algorithm to calculate the maximum cumulative likelihoods with your feature data and models, according to the methods defined in your lectures
  - b. Ensure that the computed maximum cumulative likelihoods are numerically correct and do not suffer from the underflow problem
  - c. Determine the recognition outputs, score the results and calculate the error rate
  - d. Repeat up to a total of 15 iterations or training epochs
5. Evaluate the recognizer on the supplied test data in the evaluation set
  - a. Compute the maximum cumulative likelihoods via the Viterbi algorithm
  - b. Score the recognition outputs and derive the confusion matrix
  - c. Repeat the evaluation on the team's recorded test data, reporting the score and confusions
  - d. Provide an analysis and summary of the key results using your recognizer

So there are 5 main parts of this coursework and these are divided into 18 smaller activities in total. **As a group, you will need to decide how you are going to manage and coordinate your work.** You will need to distribute activities across the team to enable each individual to make a contribution, according to their abilities. You should expect to adapt your plan, as work evolves, perhaps on a weekly basis.

For this assignment, it is expected that your calculations are coded in Matlab from scratch, using the built-in functions and libraries available to university users. You will acquire the best understanding of the algorithms working this way. Note that, if you employ third party implementations for any of the required components, such as a toolbox, toolkit, library or software downloaded from an online repository, you must highlight the lines of code that have been contributed by others and **provide the appropriate citation to the source in your references**. Any third-party code that is not appropriately cited in your report may be regarded as academic misconduct as an act of plagiarism, and will be attributed to the author of the contribution chapter where this occurs.

## 5. Assessment

Your assignment mark will take into consideration both your individual contribution and the group's overall achievements. Therefore, it is critical that each report explains in the introduction how responsibility for the tasks has been divided, stating who is the author of each contribution chapter.

Assessment of your individual contribution will be based primarily on whether the results in your individual contribution chapter are complete and correct. The group's overall achievement includes the quality of your group's technical report, which contains the abstract, introduction, planning of human resources over the project, summary of conclusions, bibliography and any appendices, in addition to those individual contribution chapters.

*PJBJ*